
Cash Is Not Zero: Auditing a Weight-Centric Quantitative Trading Pipeline

Kevin Lee 

University of California at Los Angeles
micro1
kevinlee69720@g.ucla.edu

Abstract

Quantitative trading systems can produce internally consistent backtests while misstating investment performance through omitted cash returns, unmatched benchmarks, hidden transaction-cost assumptions, or research and execution paths that implement different semantics. We present an audit-first evaluation of QUANT-CORTEX, an open-source, weight-centric research and paper-trading pipeline, and use a fixed six-ETF rotation strategy as a deliberately untuned case study. The audit formalizes residual-cash accounting, next-bar execution, mandatory costs, point-in-time contracts, coherent execution state, and machine-readable provenance. On 2,011 U.S. trading sessions from 2018–2025, residual cash earns an SHV adjusted-close proxy and the strategy is compared with passive returns scaled to the same daily risky exposure. The audited strategy earns a 1.33% net CAGR, a -0.14 cash-excess Sharpe ratio, and a -10.33% maximum drawdown. Before costs, it trails an exposure-matched equal-initial-weight buy-and-hold basket by 3.08 percentage points per year in arithmetic return; a paired circular-block bootstrap gives a 95% interval of [-5.58%, -0.56%]. All four predefined overlay variants trail their matched passive comparator, and the conclusion persists across cost levels, subperiods, and two backtest engines. The contribution is not a new alpha model; it is a reproducible negative result and a concrete protocol for distinguishing software correctness from economic evidence.

1 Introduction

Modern quantitative platforms increasingly unify data processing, feature construction, portfolio allocation, simulation, and broker execution [Yang et al., 2020, Liu et al., 2020, Yang et al., 2026]. This integration reduces friction, but it also makes accounting and interface assumptions load-bearing. A strategy that is 40% invested has a 60% cash position whether or not a backtest models it. A benchmark with 100% risky exposure does not answer whether security selection added value at 40% exposure. Likewise, agreement between a chart and a README does not establish causality, cost realism, or execution safety.

These issues are especially consequential in financial machine learning, where observations are dependent, effective sample sizes are small, and model or strategy selection creates substantial false-discovery risk [White, 2000, Harvey et al., 2016, Bailey and de Prado, 2014, Bailey et al., 2017, Arnott et al., 2019]. Reproducible software is necessary but not sufficient: the reproduced quantity must itself be economically well-defined.

This paper makes three contributions. First, it specifies an auditable return contract for weight-centric pipelines: causal target weights, explicit residual cash, mandatory costs, aligned comparators, and separable asset, cash, and cost contributions. Second, it describes engineering controls that bind reported results to code, inputs, dependencies, and execution state. Third, it reports a fixed negative-

result study of a multi-asset rotation strategy. We report all predefined variants rather than selecting a winner, and separate gross signal quality from net implementability. The purpose is to show what changes when the accounting question is posed correctly, not to claim a profitable trading rule.

2 Related work

Portfolio construction originates from explicit trade-offs between return and risk [Markowitz, 1952]; momentum evidence motivates ranking assets by past returns across both securities and asset classes [Jegadeesh and Titman, 1993, Moskowitz et al., 2012]. The case-study strategy uses these ideas only as a test vehicle. We do not claim that its particular grouping, residualization, or overlays improve on that literature.

Open platforms such as Qlib and FinRL provide modular workflows for financial AI research [Yang et al., 2020, Liu et al., 2020]; FinRL-X emphasizes a common weight interface across research and deployment [Yang et al., 2026]. Our focus is narrower and complementary: the accounting and state invariants that must hold when such an interface is evaluated or sent toward a broker.

Sharpe-ratio inference is sensitive to distributional and serial-dependence assumptions [Lo, 2002]. Data snooping and repeated trials further inflate apparent performance [White, 2000, Bailey and de Prado, 2014, Bailey et al., 2017]. We therefore avoid selecting among variants and use paired block resampling for uncertainty. More broadly, the manifest, locked environment, and executable protocol follow the reproducibility goals described by Pineau et al. [2021].

3 Audit protocol

3.1 Weight, timing, and cash contracts

Let $w_t \in \mathbb{R}^N$ be the risky-asset target decided after observing the close at date t . In the audited convention, the target executes on the first bar strictly after t and begins earning returns only after that execution bar. Thus close-derived signals cannot earn the same close they observe. For the long-only case, $w_{i,t} \geq 0$ and $\sum_i w_{i,t} \leq 1$; the residual cash weight is

$$w_t^c = 1 - \sum_{i=1}^N w_{i,t}. \quad (1)$$

For asset returns $r_{i,t}$ and cash-account return r_t^c , pre-cost portfolio return is

$$r_t^{\text{gross}} = \sum_{i=1}^N w_{i,t-1} r_{i,t} + w_{t-1}^c r_t^c. \quad (2)$$

Missing cash observations fail the run instead of silently becoming zero. The after-cost return is $r_t^{\text{net}} = r_t^{\text{gross}} - c_t$. The baseline uses a symmetric 13 basis-point charge per dollar traded with no commission or tax decomposition. This flat model is a first-order sensitivity assumption, not a market-impact model; large or urgent orders require size- and liquidity-aware execution models such as Almgren and Chriss [2001].

We report nominal CAGR, maximum drawdown, annualized one-way turnover, and cash-excess Sharpe

$$\text{SR}_c = \sqrt{252} \frac{\text{mean}(r_t - r_t^c)}{\text{sd}(r_t - r_t^c)}. \quad (3)$$

An SHV adjusted-close series is used as an investable short-duration cash proxy. It is not a frictionless risk-free asset.

3.2 System controls

The implementation separates data, alpha, portfolio, timing, risk, backtest, execution, and strategy layers under one Python package. Allocation contracts are strict at strategy output and relaxed only when an overlay intentionally reduces gross exposure. Point-in-time merge guards reject fundamentals that were not yet announced. Optional broker, machine-learning, and provider SDKs remain lazy imports so the offline scientific core is testable without credentials.

Research outputs include the generator hash, Git revision, Python and package versions, input SHA-256 digest, full experiment design, and hashes of every table and figure. Paper-trading state is stored as one versioned snapshot of positions, intents, orders, and metadata; writes are atomic and use optimistic concurrency. An order intent is persisted before submission, and an uncertain submission is not retried automatically. Broker conformance tests instantiate the current Alpaca and Interactive Brokers SDK request types, but authenticated transport and venue behavior remain outside this study.

4 Case study and experimental design

4.1 Fixed strategy

The universe contains QQQ, VGT, GLD, TLT, SPY, and VIG, grouped as growth, real assets, and defensive proxies. On the first session of each week, the strategy ranks groups by a 126-session information ratio relative to QQQ and keeps the top two. Within selected groups, it estimates a one-factor regression on a preceding 126-session window and sums residual returns over a separate 126-session formation window, skipping the latest 21 sessions. Positive scores receive proportional long-only weights, capped at 60% per asset. QQQ is also one of the two growth-group members, so that group's active-return spread contains one zero benchmark leg; this asymmetric product specification is held fixed.

Two exposure overlays are evaluated. A seeded three-state Gaussian-mixture regime model maps QQQ return and realized-volatility features to zero, half, or full exposure. A volatility scaler targets 20 annualized volatility points; no VIX series is supplied, so the experiment uses the documented 21-session QQQ realized-volatility proxy. The full strategy uses both overlays. Three predefined variants remove the regime gate, remove the volatility scaler, or remove both. This audit did not tune parameters or select a variant based on the reported period. The unknown number of configurations explored before the audit prevents a defensible multiple-testing-adjusted discovery claim.

4.2 Data, comparators, and uncertainty

The owner-supplied matrix contains adjusted closes from 2016-01-04 through 2025-12-31, retrieved on 2026-06-16 through yfinance 1.4.1 with `auto_adjust=False`. The yfinance documentation describes the Yahoo Finance API as intended for personal use and directs users to Yahoo's terms of use.¹ The repository owner attests permission to publish the derived aggregate results; the raw matrix is not distributed. Its SHA-256 is:

efb384a62157e56a0cd8065abf45c1ed07d90ec26c681e5d54d74fe4cb9c55e1

The first 503 sessions are signal history; evaluation covers 2,011 sessions from 2018-01-02 through 2025-12-31.

Comparators share one capital clock. SPY and a six-ETF basket initialized at equal dollar weights are measured from the close immediately preceding the evaluation window. To isolate cross-sectional decisions from exposure timing, the basket return is multiplied by the strategy's realized risky exposure each day and the remainder earns SHV. This comparator is gross of its own trading costs. We therefore report both gross-to-gross signal comparisons and the more conservative net-strategy-to-gross-benchmark comparison.

For dependent daily observations, we resample paired active-return series with the circular-block bootstrap of Politis and Romano [1992]. We use 5,000 replications, 21-session blocks, and seed 42. Intervals quantify sampling variation conditional on this one historical path, strategy specification, cash proxy, and block length; they do not establish stationarity or future performance.

¹<https://ranaroussi.github.io/yfinance/>

Table 1: Audited accounting over 2018–2025. Sharpe ratios use SHV except for the explicitly zero-cash row, which uses a zero reference rate.

Series	CAGR	Cash-excess Sharpe	Max drawdown
Strategy, after costs and SHV cash	1.33%	-0.14	-10.33%
Strategy, before costs and SHV cash	3.26%	0.15	-8.61%
SHV cash proxy	2.50%	–	-0.45%
Exposure-matched passive basket	6.54%	0.69	-7.97%
Strategy, after costs and zero cash	-0.07%	0.02	-14.22%




Figure 1: Left: growth of one dollar under audited cash and cost accounting. Right: calendar-year nominal returns. The matched comparator uses the strategy’s daily risky exposure and assigns residual capital to SHV.

5 Results

5.1 Cash accounting changes the reported level, not the verdict

Table 1 and Figure 1 separate cash and cost effects. Crediting residual cash raises net CAGR from -0.07% under a zero-return cash assumption to 1.33%. That correction is economically material because mean risky exposure is only 34.76% and the strategy is fully in cash on 41.92% of sessions. It does not create evidence of alpha: SHV alone compounds at 2.50%, and the exposure-matched passive basket compounds at 6.54%.

5.2 Costs and overlays do not rescue the strategy

The strategy turns over 10.84 times per year on a one-way basis. At zero cost, its cash-excess Sharpe is 0.15; it declines to 0.03 at 5 bps, -0.14 at the 13-bps baseline, -0.41 at 25 bps, and -0.94 at 50 bps. The model is therefore not economically robust to modest execution friction.

Figure 2 compares each variant before costs with a passive basket at the same exposure. This avoids attributing the result solely to strategy turnover. The full model has gross cash-excess Sharpe 0.15 versus 0.69 for its comparator. Removing the regime gate, the volatility scaler, or both increases strategy exposure and raw returns, but every gross strategy still trails its matched comparator. The corresponding annualized gross active returns are -3.08%, -4.63%, -3.21%, and -4.28%. No variant is selected as a replacement model.

5.3 Uncertainty, subperiods, and engine parity

The net strategy’s annualized arithmetic excess return over SHV is -0.94%, with a 95% block-bootstrap interval of [-5.27%, 3.55%]. Before costs, the corresponding estimate is 0.95% [-3.31%, 5.40%]. Neither supports a positive cash-relative claim. Relative to the exposure-matched basket, gross active return is -3.08% [-5.58%, -0.56%], and only 0.64% of bootstrap replications are positive. After costs it is -4.97% [-7.58%, -2.39%], with no positive replication among 5,000 draws.




Figure 2: Left: after-cost Sharpe sensitivity for the full strategy. Right: gross-to-gross ablations, so transaction costs do not explain the passive comparator’s advantage. All variants were defined before this comparison.

The conclusion also holds in both prespecified halves. In 2018–2021, strategy CAGR is 0.08% and cash-excess Sharpe is -0.18, versus 4.77% and 0.78 for the matched basket. In 2022–2025, the values are 2.59% and -0.12, versus 8.34% and 0.66. Finally, the vectorized and share-based event-driven engines produce net CAGRs of 1.325% and 1.328%, respectively, and cash-excess Sharpes of -0.1442 and -0.1440. This close agreement is a diagnostic for this experiment, not a claim that their rebalancing semantics are generally equivalent.

6 Limitations and broader impact

This study evaluates one U.S.-listed ETF universe over one eight-year window. SHV has duration, expenses, and market-price variation, so it is only a cash proxy. Adjusted closes omit intraday spreads and fill uncertainty. The flat cost grid excludes order-size, volatility, venue, and market-impact dynamics. The exposure-matched basket is a diagnostic comparator, not a directly quoted security, and is reported gross of its own rebalancing cost. The bootstrap is conditional on a weak-dependence approximation and a chosen 21-session block. Historical research trial counts are unknown, so a Deflated Sharpe Ratio would require an assumed rather than observed number of trials. The data source can revise adjusted histories; the exact input is identified by digest but is not openly redistributed. No live broker authentication, exchange acknowledgment, or funded paper order was tested.

The positive impact of an audit-first protocol is reduced false confidence: cash, costs, timing, and negative results become visible rather than optimized away. The principal risk is the opposite misuse: readers may treat an audited backtest as investment advice or deployment approval. The repository therefore defaults to paper mode, requires pre-trade risk checks, blocks uncertain order retries, and describes the strategy as a research baseline. Past simulated performance does not imply future returns.

7 Conclusion

Software correctness and investment efficacy are different hypotheses. In this case, explicit cash accounting materially improves the reported CAGR, while matched-exposure benchmarks, gross-to-gross ablations, and dependence-aware uncertainty show that the fixed strategy does not add value over simple passive exposure. The broader lesson is procedural: weight-centric trading research should make cash a first-class asset, separate gross signal from net implementation, compare like exposure with like exposure, preserve every predefined variant, and bind artifacts to code and inputs. A negative result under those controls is more useful than a stronger chart built on ambiguous accounting.

Acknowledgments

No external funding was used for this work.

References

- Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *The Journal of Risk*, 3(2):5–39, 2001. doi: 10.21314/JOR.2001.041.
- Rob Arnott, Campbell R. Harvey, and Harry Markowitz. A backtesting protocol in the era of machine learning. *The Journal of Financial Data Science*, 1(1):64–74, 2019. doi: 10.3905/jfds.2019.1.064.
- David H. Bailey and Marcos López de Prado. The deflated sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *The Journal of Portfolio Management*, 40(5):94–107, 2014. doi: 10.3905/jpm.2014.40.5.094.
- David H. Bailey, Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu. The probability of backtest overfitting. *The Journal of Computational Finance*, 20(4):39–69, 2017. doi: 10.21314/JCF.2016.322.
- Campbell R. Harvey, Yan Liu, and Heqing Zhu. ... and the cross-section of expected returns. *The Review of Financial Studies*, 29(1):5–68, 2016. doi: 10.1093/rfs/hhv059.
- Narasimhan Jegadeesh and Sheridan Titman. Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1):65–91, 1993. doi: 10.1111/j.1540-6261.1993.tb04702.x.
- Xiao-Yang Liu, Hongyang Yang, Qian Chen, Runjia Zhang, Liuqing Yang, Bowen Xiao, and Christina Dan Wang. FinRL: A deep reinforcement learning library for automated stock trading in quantitative finance. *arXiv preprint arXiv:2011.09607*, 2020. URL <https://arxiv.org/abs/2011.09607>.
- Andrew W. Lo. The statistics of sharpe ratios. *Financial Analysts Journal*, 58(4):36–52, 2002. doi: 10.2469/faj.v58.n4.2453.
- Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952. doi: 10.1111/j.1540-6261.1952.tb01525.x.
- Tobias J. Moskowitz, Yao Hua Ooi, and Lasse Heje Pedersen. Time series momentum. *Journal of Financial Economics*, 104(2):228–250, 2012. doi: 10.1016/j.jfineco.2011.11.003.
- Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alché Buc, Emily Fox, and Hugo Larochelle. Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, 22(164):1–20, 2021. URL <https://jmlr.org/papers/v22/20-303.html>.
- Dimitris N. Politis and Joseph P. Romano. A circular block-resampling procedure for stationary data. In Raoul LePage and Lynne Billard, editors, *Exploring the Limits of Bootstrap*, pages 263–270. John Wiley & Sons, New York, 1992.
- Halbert White. A reality check for data snooping. *Econometrica*, 68(5):1097–1126, 2000. doi: 10.1111/1468-0262.00152.
- Hongyang Yang, Boyu Zhang, Yang She, Xinyu Liao, and Xiaoli Zhang. FinRL-X: An AI-native modular infrastructure for quantitative trading. *arXiv preprint arXiv:2603.21330*, 2026. URL <https://arxiv.org/abs/2603.21330>.
- Xiao Yang, Weiqing Liu, Dong Zhou, Jiang Bian, and Tie-Yan Liu. Qlib: An AI-oriented quantitative investment platform. *arXiv preprint arXiv:2009.11189*, 2020. URL <https://arxiv.org/abs/2009.11189>.

A Reproducibility details

The complete generator is `scripts/run_paper_experiments.py`. From a locked development environment, the fixed command is:

```
PYTHONPATH=. .venv/bin/python scripts/run_paper_experiments.py \
  --prices-csv local_data/published_rotation_prices.csv \
  --output-dir paper --bootstrap-replications 5000
```

The input must contain adjusted-close columns `QQQ`, `VGT`, `GLD`, `TLT`, `SPY`, `VIG`, and `SHV` and match the published digest for exact reproduction. Aggregate CSVs, figures, and their hashes are in `paper/results/manifest.json`. The run used Python 3.14.4 on an eight-core Apple M1 with 8 GiB RAM and completed in approximately 25 seconds on CPU; no GPU was used. Supported CI runs use Python 3.11 through 3.14. The reported experiment consumed less than one CPU-minute; repository-wide testing, packaging, notebook execution, and broker conformance checks were additional engineering validation rather than model search.

B Additional engine comparison




Figure 3: The vectorized engine re-pegs target weights between rebalances, whereas the event-driven engine holds explicit drifting shares. Their small difference here does not remove the semantic distinction.

C Use of language-model assistance

An LLM-based coding assistant supported repository inspection, implementation review, regression-test generation, and language editing. It was not a component of the strategy, data, statistical method, or experiment-selection rule. Numerical claims in the paper are generated from committed scripts and machine-readable artifacts; responsibility for the work remains with the author.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Introduction state that this is an audit protocol and negative-result case study, not a new alpha model or deployment claim. The numerical claims match Sections 4 and 5.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 discusses the single-universe design, cash proxy, adjusted-close data, flat costs, comparator construction, bootstrap assumptions, unknown prior trial count, data availability, and untested live transport.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains accounting definitions and empirical estimators, but no theorem, lemma, or new theoretical result requiring a proof.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: Sections 3–4 and Appendix A disclose the algorithms, dates, symbols, parameters, versions, command, and exact input digest, but the raw price matrix is not redistributed, so bitwise independent reproduction is not fully open.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The MIT-licensed code, aggregate results, figures, and manifests are open, but the provider-derived adjusted-close matrix is not. Appendix A explains the exact command and required columns for an authorized local copy.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 specifies the full fixed strategy, lookbacks, execution timing, variants, cost grid, evaluation split, comparators, bootstrap settings, and the absence of selection or tuning in this audit.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 5 reports paired circular-block bootstrap intervals and positive-replication frequencies, while Section 4 states the 21-session block, 5,000 replications, seed, dependence rationale, and conditional interpretation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A reports CPU model, core count, memory, Python version, runtime, GPU use, and distinguishes experiment compute from broader engineering validation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses public financial instruments, no human subjects, no private personal data, and no live trading. Section 6 describes release safeguards and misuse risks.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 6 discusses the benefit of reducing false investment confidence and the risks of treating simulated results as advice or deployment approval, together with practical mitigations.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no pretrained model, scraped media corpus, personal data, or other high-risk dual-use asset requiring gated access.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Section 4 identifies Yahoo Finance via yfinance 1.4.1, links its usage notice, records the owner's permission attestation for derived aggregates, and states that raw quotes are not redistributed. Related methods are cited in Sections 2–4.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The repository documents the MIT-licensed code, architecture, tests, security policy, experiment command, aggregate result tables, figures, provenance manifest, and known limitations.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research includes no crowdsourcing and no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The research includes no intervention or data collection involving human subjects, so IRB review is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not an important, original, or non-standard component of the strategy or empirical method. Appendix C voluntarily discloses coding and language-editing assistance.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.